



# Systems Programming

## Lecture 7:

### Inter Process Communication: Semaphores and Messages

Dmitry A. Zaitsev

<https://dimazaitsev.github.io>

# Semaphore

- E.W.Dijkstra, 1962
- Integer variable S
- Initialization  $S = <\text{number of resources}>$
- Operations: P – wait, V – signal
- $P(S)$ : if( $S > 0$ )  $S = S - 1$ ; else <wait on S>
- $V(S)$ : if(<there are waiting on S>) <wake up first waiting>; else  $S = S + 1$ ;

# Producers - Consumers

emptyCount=N; fullCount=0; useQueue=1;

- Produce:

P(emptyCount)

P(useQueue)

putItemIntoQueue(Item)

V(useQueue)

V(fullCount)

- Consume:

P(fullCount)

P(useQueue)

Item =getItemFromQueue()

V(useQueue)

V(emptyCount)

# Linux POSIX semaphores

- `#include <fcntl.h> /* For O_*` constants \*/  
`#include <sys/stat.h> /* For mode constants */`  
`#include <semaphore.h>`
- Link -lpthread
- Operations:
  - P(S): `sem_wait(s)`
  - V(S): `sem_post(s)`
- Name – external, pointer to *sem\_t* - internal

# Initialize and open a named semaphore

- `sem_t *sem_open(const char *name, int oflag);`
- `sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value);`
- The *oflag* argument specifies flags that control the operation of the call: `O_CREAT`, `O_EXCL`
- The *mode* argument specifies the permissions to be placed on the new semaphore.
- The *value* argument specifies the initial value for the new semaphore.

# Unlock and Lock a semaphore

- Unlock a semaphore

```
int sem_post(sem_t *sem);
```

- Lock a semaphore

```
int sem_wait(sem_t *sem);
```

```
int sem_trywait(sem_t *sem);
```

```
int sem_timedwait(sem_t *sem, const struct  
timespec *abs_timeout);
```

# Finish using semaphore

- close a named semaphore

```
int sem_close(sem_t *sem);
```

- remove a named semaphore

```
int sem_unlink(const char * name);
```

# More functions for semaphores

- get the value of a semaphore

```
int sem_getvalue(sem_t *sem, int *sval);
```

- initialize an unnamed semaphore

```
int sem_init(sem_t *sem, int pshared, unsigned  
int value);
```

- destroy an unnamed semaphore

```
int sem_destroy(sem_t *sem);
```

# A task: dining philosophers

- Pension
- 5 philosophers: thinking or eating
- 5 plates and 5 forks among them
- A philosopher needs 2 forks to start eating
- Deadlock !
- How to avoid deadlocks ?
- Timeouts and other solutions

# Messages

- A queue in kernel
- A process sends a message to a queue
- A message is stored in a queue
- If message queue is full, block the sender
- A process receives a message from queue
- If message queue is empty, block the receiver

# POSIX message queues

- #include <mqueue.h>
- Link -lrt
- Name: “/somename”
- Descriptor mqd\_t
- Attributes

```
struct mq_attr {  
    long mq_flags;      /* Flags: 0 or O_NONBLOCK */  
    long mq_maxmsg;    /* Max. # of messages on queue */  
    long mq_msgsize;   /* Max. message size (bytes) */  
    long mq_curmsgs;   /* # of messages currently in queue */  
};
```

- Basic operations: mq\_send(...), mq\_receive(...)

# Open a message queue

- `mqd_t mq_open(const char *name, int oflag);`
- `mqd_t mq_open(const char *name, int oflag, mode_t mode, struct mq_attr *attr);`
- The *oflag* argument specifies flags that control the operation of the call: `O_RDWR`, `O_CREAT`, `O_NONBLOCK`
- The *mode* argument specifies the permissions to be placed on the new queue
- The *attr* argument specifies attributes for the queue.

# Send a message

- `mq_send(mqd_t mqdes, const char *msg_ptr,  
size_t msg_len, unsigned int msg_prio);`
- `int mq_timedsend(mqd_t mqdes, const char  
*msg_ptr, size_t msg_len, unsigned int  
msg_prio, const struct timespec  
*abs_timeout);`
- `struct timespec {  
 time_t tv_sec; /* seconds */  
 long tv_nsec; /* nanoseconds */  
};`

# Receive a message

- `ssize_t mq_receive(mqd_t mqdes, char *msg_ptr, size_t msg_len, unsigned int *msg_prio);`
- `ssize_t mq_timedreceive(mqd_t mqdes, char *msg_ptr, size_t msg_len, unsigned int *msg_prio, const struct timespec *abs_timeout);`

# More functions for message queues

- `int mq_close(mqd_t mqdes);`
- `int mq_unlink(const char *name);`
- `int mq_getattr(mqd_t mqdes, struct mq_attr *attr);`
- `int mq_setattr(mqd_t mqdes, const struct mq_attr *newattr, struct mq_attr *oldattr);`

```
#include "mq_common.h"
int main(int argc, char **argv)
{
    mqd_t mq;
    struct mq_attr attr;
    char buffer[MAX_SIZE + 1];
    int must_stop = 0;
    attr.mq_flags = 0;
    attr.mq_maxmsg = 10;
    attr.mq_msgsize = MAX_SIZE;
    attr.mq_curmsgs = 0;
    mq = mq_open(QUEUE_NAME, O_CREAT | O_RDONLY, 0644, &attr);
    CHECK((mqd_t)-1 != mq);
    do { ssize_t bytes_read;
        bytes_read = mq_receive(mq, buffer, MAX_SIZE, NULL);
        CHECK(bytes_read >= 0);
        buffer[bytes_read] = '\0';
        if (!strncmp(buffer, MSG_STOP, strlen(MSG_STOP))) must_stop = 1;
        else printf("Received: %s\n", buffer);
    } while (!must_stop);
    CHECK((mqd_t)-1 != mq_close(mq));
    CHECK((mqd_t)-1 != mq_unlink(QUEUE_NAME));
    return 0;
}
```

## Messages: server

# Messages: client

```
#include "mq_common.h"
int main(int argc, char **argv)
{
    mqd_t mq;
    char buffer[MAX_SIZE];
    mq = mq_open(QUEUE_NAME, O_WRONLY);
    CHECK((mqd_t)-1 != mq);
    printf("Send to server (enter \"exit\" to stop it):\n");
    do {
        printf("> ");
        fflush(stdout);
        memset(buffer, 0, MAX_SIZE);
        fgets(buffer, MAX_SIZE, stdin);
        CHECK(0 <= mq_send(mq, buffer, MAX_SIZE, 0));
    } while (strncmp(buffer, MSG_STOP, strlen(MSG_STOP)));
    CHECK((mqd_t)-1 != mq_close(mq));
    return 0;
}
```

# Messages: common

```
#ifndef COMMON_H_
#define COMMON_H_

#define QUEUE_NAME "/test_queue"
#define MAX_SIZE 1024
#define MSG_STOP "exit"

#define CHECK(x) \
do { \
    if (!(x)) { \
        fprintf(stderr, "%s:%d: ", __func__, __LINE__); \
        perror(#x); \
        exit(-1); \
    } \
} while (0)

#endif /* #ifndef COMMON_H_ */
```